

NoSQL Databases

Polytech Nice-Sophia, SI4

Pierre Monnin

pierre.monnin@inria.fr

Chapter 2:

Document Databases

Overview

- **Key idea:**

- Store/Load/Modify documents with arbitrary fields in a database*

- Simple text search functionality*

- **Difference from Column Storage:**

- You do not need to define any schema – it is defined by the data you insert*

Schema (1/2)

- Documents bring their own schema – they are self-describing
 - Formats: XML, JSON etc.*
- No need to define data types – detected automatically
- Indexes are possible
- Documents are part of collections/databases
- Documents are versioned
 - Keep in mind to clean them up regularly*

Schema (2/2)

- References („foreign keys“) are managed by the application
- Documents can share parts of their schema, but none of it as well
- Do not let documents grow too large to avoid swapping to disk
 - *normalization makes sense*
 - *think about what your application really needs*

Schema: Example /Json

-Document 1

```
-{  
  „username“: „maxmustermann“,  
  „likes“: [„soccer“, „photography“]  
-}
```

-Document 2

```
-{  
  „username“: „marthamusterfrau“,  
  „profession“: „CEO“,  
  „businesstrips“: [„Bangkok“, „New York“, „Brisbane“]  
  „address“: {  
    street: „1569 Broadway“,  
    city: „New York“  
  }  
-}
```

Type of Queries (1/2)

- Different from SQL
- Depends on document store
- Usually document-oriented
 - Creating, Modifying, Retrieval (Simple Search)*

Type of Queries (2/2)

- Simple aggregations
- Java/Javascript syntax
- Additional functionality to implement any kind of query
 - Map/Reduce Paradigm*

Distributed Persistence

- **Replication**

- Usually Master/Slave paradigm

- Arbiter: elects new master in case of failure

- You can define read/write-quorums for each command you execute

- **Sharding**

- By a key in a document

- Range-Sharding: define a range which should be stored on a node

- Example: All documents from 01.01.2014 – 31.12.2014

- Good for data processing on specific nodes

- Hash-Sharding: ensures that data is evenly distributed across nodes

- Good for heavy write operations

Transactional Guarantees

- Operations on a single document are always atomic (MongoDB)
- Update document only if version number in update is the same as version number in document (CouchDB)

*–Example: Update 1 {Version: 1, title: „this is a test“} can be applied
Document {Version: 1, title „this is nothing“}*

–Similar to vector clock approach

- Multiple documents:

*–Create your own transaction mechanism using 2PC (2 Phase Commit)
protocol*

*–Alternatively: Have one document with nested documents (→ not always
recommended)*

Technologies

- **MongoDB**

- Very popular database*

- Master/Slave Replication*

- GNU AGPL V3.0*

- **CouchDB**

- Similar functionality as MongoDB*

- Master/Master replication*

- Apache License*

Use Cases

- Web-Analytics
 - Store all events on the web page (clicks, scroll, touch, ...) - usually already JSON-objects*
- Catalogue/Shopping applications
 - Products with different fields*
- Reviews / Tweets

Evaluation

- Document stores are very useful where a flexible schema is required
- Good for prototyping – no need to think much about schema beforehand
- Fixed schemas / data types can offer more performance for specific use cases
- Relations between documents have to be managed by the application

Document Databases -

MongoDB Exercises

Overview

- MongoDB was initially released in 2009
- Open Source (GNU AGPL v3.0)
- Name comes from **humongous**
- Typical document storage
- Schema is defined by the data you insert
- JavaScript is used to create more complex queries

Schema

- Collections (~ databases): Contain documents
- Documents
 - *Internal representation: JSON Format*
 - *Various formats can be imported*
- Further schema details are defined by the document you insert

Excursus: JSON-Format

- JavaScript Object Notation
- Relevant for NoSQL and Web applications
- You define data types implicitly
- Encoding: Often UTF-8
- Example

```
{  
  "name": "Homer",  
  "isPoet": true,  
  "age": 35,  
  "height": 180.1,  
  "address": {  
    "city": "Troy"  
  "country": "Turkey"  
}
```

Queries

- Syntax similar to Javascript
- Aggregation calculation support several different type of aggregations for analytics
- Geospatial Queries
 - Example: Give me all tweets near my location*
- You can use MapReduce for advanced calculations

Queries: CRUD

- CRUD = Create Read Update Delete

-Common term for basic data operations

- **Create**

-db.collection.insert(document)

-Example: db.test.insert({_id:10, type: "test"})

- **Read** (can match multiple documents!)

-db.collection.find(criteria)

-Example: db.collection.find({_id:10})

- **Update** (can match multiple documents!)

-db.collection.update(criteria, update, options)

-Example: db.collection.update({_id:10}, {type: "final"}, {multi:false})

- **Delete**

-db.collection.remove(criteria) (can remove multiple documents!)

-Example: db.collection.remove({_id:10})

Query criteria

- You can match fields exactly in a document

-Example `({type: "test"})` => all documents with the key „type“ and the corresponding value „test“

- You can use wildcards/regular expressions

-Example `({type:{$regex: "te", $options: "i"})` => all documents having a type starting with te*

- You can use `> ($gt)` or `< ($lt)` for numbers

-Example `({price:{$gt:100})`

- You can use OR conditions (default AND)

-Example `({$or: [{type: "test"}, {price:{$gt:100}}]})`

Queries: Aggregation

- **Simple**

– *Count: Count all documents matching criteria*

• *Example: db.collection.find(criteria).count()*

– *Distinct: Returns all unique value for a given field*

• *Example: db.collection.distinct(“type”)*

– *Group (Warning: Not usable for sharded collections, limit of the documents returned)*

- **Aggregation Pipeline**

– *You can define a pipeline of several aggregation commands (e.g. sum, avg, first, last, max, min, ...)*

– *MongoDB optimizes execution*


- **MapReduce**

– *Full flexibility, but more complex*

Transactions

- Only on document level
- Not across documents or collections
- As always: You can develop any transaction logic on top

Architecture



Distributed Persistence

- **Replication**

- Client writes to primary

- Primary replicates to secondary

- Client may read from secondaries for load balancing

- .May have outdated data

- Primary fails → Config Servers support Secondary to elect a new primary among them

- **Sharding**

- Query router routes the queries to the correct shard

- You define an attribute of a document as a shard key

- .Hash-based shard vs. range-based shard

Interfaces

- JDBC
- Libraries for various programming languages
- mongo – shell program

Excursus: Loading data into mongodb

- Mongo-importer (for documents in JSON/CSV)
- ETL Tools
 - Pentaho Kettle/Spoon*
- Write your own import tool

Who uses it?

- LinkedIn: Learning Platform
- Expedia: Travel planning
- SAP: Enterprise Content Management
- MTV: Content Management System
- Springer: Real-time analytics for downloads of journals / book chapters
- Github: Reporting
- Nearley: user interaction and activities
- O2: Several applications
- E-Bay: Various Services