

Architecture Seq2Seq et attention

Michel.RIVEILL@univ-cotedazur.fr

2024-2025

Sommaire


- ▶ **Architecture Seq2Seq**
 - ▶ Construction
 - ▶ Entrainement
 - ▶ Inférence
- ▶ **Le mécanisme d'attention**
- ▶ **Les attentions en Keras**

Vecteurs d'embedding

- ▶ J'ai bien aimé vos propositions
 - ▶ Paella + France – Espagne = Cassoulet
 - ▶ Pasta + Japon - France = noodle
 - ▶ Steak + Vegetable – Meat = Salad
 - ▶ Pancake + France – Canada = Croissant
 - ▶ Shawarma + Turkish – Lebanon = Kebab
 - ▶ Rome + Romania – Italie = Bucharest
 - ▶ Paris + Dutch – Berlin = French
 - ▶ Germany + Paris – France = Berlin
 - ▶ Boy + mother – father = girl
 - ▶ Planet + dry – water = martian
 - ▶ Music + jazz – blues = classical
- ▶ J'espère que cela vous aura permis de bien comprendre
'l'importance d'un bon embedding'


LSTM cell

- ▶ Cellule composée de trois « portes » : ce sont des zones de calcul qui régulent le flux d'informations :
 - ▶ Forget gate (porte d'oubli)
 - ▶ Input gate (porte d'entrée)
 - ▶ Output gate (porte de sortie)
- ▶ Cell state (état de la cellule)
 - ▶ Même fonction que 'residual'
 - ▶ Code une sorte d'agrégation des données de tous les pas de temps précédents qui ont été traités
- ▶ Hidden state (état caché)
 - ▶ N'a rien de cache... puisque c'est celui qui l'on observe
 - ▶ Code une sorte de caractérisation des données du pas de temps précédent.



GRU – gated recurrent unit

- GRU est plus simple que le LSTM.



$$z_t = \sigma (W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma (W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh (W \cdot [r_t * h_{t-1}, x_t])$$


$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

- GRU combine l'oubli et l'entrée en une seule porte de mise à jour.
- GRU fusionne également l'état de la cellule et l'état caché.


Architecture Seq2Seq

Auto-encodeur avec cellule récurrente


Architecture des auto-encodeurs



Exemple d'utilisation des auto-encodeurs




Auto-encodeurs avec des CNN



Auto-encoder avec des cellules récurrentes

- ▶ Étendre l'architecture du codeur-décodeur
 - ▶ à une séquence de données
 - ▶ afin de développer une architecture capable de générer des séquences de sortie de longueur arbitraire adaptées au contexte



Seq2Seq model

- ▶ Applications with text
 - ▶ Machine translation
 - ▶ Text summarization
 - ▶ Question answering
 - ▶ Dialogue modeling
- ▶ But also : forecasting
- ▶ Il est aussi possible de remplacer l'encoder par un autre latent space :
 - ▶ Par exemple : Image captioning






Image captioning model



Seq2seq – original approach

- Extend encoder-decoder architecture to a sequence data
arbitrary length




Sequence-to Sequence Architectures

- ▶ Three main part:
 - ▶ Encoder: processes the input sequence
 - ▶ Ordinary sequence-to-vector RNN
 - ▶ Context: output of the encoder
 - ▶ is usually a simple function of its final hidden state ($h + c$)
 - ▶ aims to encapsulate the information for all input elements in order to help the decoder make accurate predictions
 - ▶ Decoder: is conditioned on the context to generate the output sequence
 - ▶ the context acts as the initial hidden state of the decoder part of the model
 - ▶ produces output at each step


Comment construire une architecture Seq2Seq

- ▶ Encodeur
- ▶ Context
- ▶ Décodeur : 2 approches possibles
 - ▶ Sans teacher : plus difficile à construire et plus difficile à faire converger
 - ▶ Avec teacher : approche aujourd'hui adoptée lors du TD

Seq2Seq sans teacher



Seq2Seq avec teacher



model.fit ([x, Teacher], y)

Le 'teacher' est construit à partir de y

Pros and Cons of Teacher Forcing

▶ Pros:

- ▶ If we do not use Teacher Forcing
 - ▶ the hidden states of the model will be updated by a sequence of wrong predictions
 - ▶ **errors will accumulate**
 - ▶ and it is difficult for the model to learn from that.
- ▶ **Training with Teacher Forcing converges faster.**

▶ Cons:

- ▶ **Unfortunately, during inference, there is no ground truth available**
 - ▶ the RNN model will have to re-inject its own prediction for the next prediction.
- ▶ There is a difference between
 - ▶ learning (no propagation of error)
 - ▶ inference (propagation of error),which leads to poor performance and model instability.
- ▶ This phenomenon is known as **exposure bias** in the literature.

Comment construire un modèle Seq2Seq avec ‘teacher’

▶ Encoder

```
def build_encoder():

    inputs = layers.Input(shape=_INPUT_LENGTH,), name="enInput")

    h = layers.Embedding(input_dim=_VOCAB_SIZE,
                          output_dim=_EMBEDDING_SIZE)(inputs)

    encoder_outputs, memory_state, carry_state = layers.LSTM(_RNN_SIZE,
                                                               return_sequences=True, return_state=True,
                                                               dropout=_DROPOUT_RATE, recurrent_dropout=_DROPOUT_RATE)(h)

    encoder_context = [memory_state, carry_state]

    return models.Model(inputs, [encoder_outputs, encoder_context], name="encoder")
```

Comment construire un modèle Seq2Seq avec ‘teacher’

► Decoder

```
def build_decoder():

    teacher_inputs = layers.Input(shape=_OUTPUT_LENGTH,), name="decInput")
    ctx_input = [layers.Input(shape=_RNN_SIZE,), name="ctxInputH"),
                layers.Input(shape=_RNN_SIZE,), name="ctxInputC")]

    h = layers.Embedding(input_dim=_VOCAB_SIZE,
                          output_dim=_EMBEDDING_SIZE)(teacher_inputs)

    decoder_outputs, memory_state, carry_state = layers.LSTM(_RNN_SIZE,
                                                               return_sequences=True, return_state=True,
                                                               dropout=_DROPOUT_RATE,
                                                               recurrent_dropout=_DROPOUT_RATE)(h, initial_state=ctx_input)

    context = [memory_state, carry_state]

    outputs = layers.Dense(_VOCAB_SIZE, activation='softmax')(decoder_outputs)

    return models.Model([teacher_inputs, context_input], [outputs, context], name="decoder")
```

Comment construire un modèle Seq2Seq avec ‘teacher’

- ▶ Seq2Seq = Encoder+Decoder


```
def build_seq2seq(encoder, decoder):  
    enc_inputs = layers.Input(shape=(_INPUT_LENGTH,), name="encInput")  
    _, enc_context = encoder(enc_inputs)  
  
    teacher_inputs = layers.Input(shape=(_OUTPUT_LENGTH,), name="teacherInput")  
    dec_sequence_outputs, _ = decoder([teacher_inputs, enc_context])  
  
    return models.Model([enc_inputs, teacher_inputs], [dec_sequence_outputs])
```

Comment prédire avec une architecture Seq2Seq

- ▶ On est obligé de procéder par étape (approche gourmande)
 - ▶ On ne peut pas utiliser de teacher... car on ne connaît pas la phrase à décoder
- ▶ Le mode d'emploi est donc le suivant :
 1. Encoder la phrase
 2. Itération pour décoder successivement chaque pas de temps, en réutilisant les pas de temps déjà décodés
 3. Arrêt lorsque <fin> ou max_length
- ▶ Concrètement, il y a 3 manières de faire
 - ▶ De la plus simple et la plus inefficace
 - ▶ On fait plusieurs fois la même chose : encodage + decodage à nouveau de la partie déjà décodée
 - ▶ A la plus efficace
 - ▶ On fait une seule fois chaque étape : encodage / décodage du prochain mot en utilisant le mot que l'on vient de décoder
 - ▶ Détails dans les transparents suivants
 - ▶ En passant
 - ▶ On fait une fois l'encodage, mais on décode plusieurs fois la séquence en utilisant chaque fois, l'ensemble des étapes précédentes
- ▶ ▶ A approfondir lors du TD

Predict with seq2seq model

Step 1 : Use encode to define context




Predict with seq2seq model

Step 2 : decode the first output

Decode step by step

Reuse at each step, the previous output




DECODER RNN

Predict with seq2seq model

Step 2 : step by step... decode sentence

Decode step by step

Reuse at each step, the previous output




Predict with seq2seq model

Step 2 : step by step... decode sentence

Decode step by step

Reuse at each step, the previous output

Stop, when generate "stop" label or
max_length



DECODER RNN




On améliore la proposition Seq2Seq



Ajout du mécanisme d'attention

Sequence-to-Sequence (seq2seq)

It's crazy that the entire "meaning" of sentence is expected to be packed into this **one embedding**, and that the encoder then never interacts w/ the decoder again.




Les limites de l'approche Seq2Seq

- ▶ L'encodeur lit la phrase entière **une seule fois**
 - ▶ En produit une représentation abstraite (un vecteur)
 - ▶ Il doit donc se souvenir de toutes les parties de la phrase
 - ▶ Le décodeur utilise cette représentation pour convertir cette phrase
- ▶ Pour des phrases longues
 - ▶ le codeur ne peut pas se souvenir de la première partie de la phrase, ce qui entraînera une perte d'informations.
- ▶ Comment traduisez-vous un texte ?
 - ▶ Pensez-vous que l'ensemble de la séquence d'entrée (ou de la phrase) est important à chaque étape de l'encodage ?
 - ▶ Pouvons-nous mettre l'accent sur certains mots plutôt que d'accorder la même importance à tous les mots ?
 - ▶ Nous, les humains, on traduit chaque mot en nous concentrant principalement sur certains mots de l'entrée.
 - ▶ Idéalement, à chaque étape, on ne devrait transmettre que les informations pertinentes
 - ▶ Encodages des informations pertinentes
 - ▶ Transmission au décodeur pour la traduction
- ▶ Le mécanisme d'attention a été développé pour répondre à ces défis.


seq2seq + Attention

Q: How do we determine how much to pay attention to each of the encoder's hidden layers?




seq2seq + Attention

Q: How do we determine how much to pay attention to each of the encoder's hidden layers?
A: Let's base it on our decoder's current hidden state (our current representation of meaning) and all of the encoder's hidden layers!




seq2seq + Attention

Q: How do we determine how much to pay attention to each of the encoder's hidden layers?
A: Let's base it on our decoder's current hidden state (our current representation of meaning) and all of the encoder's hidden layers!




seq2seq + Attention

Q: How do we determine how much to pay attention to each of the encoder's hidden layers?
A: Let's base it on our decoder's current hidden state (our current representation of meaning) and all of the encoder's hidden layers!



seq2seq + Attention


Q: How do we determine how much to pay attention to each of the encoder's hidden layers?
A: Let's base it on our decoder's current hidden state (our current representation of meaning) and all of the encoder's hidden layers!



seq2seq + Attention

Q: How do we determine how much to pay attention to each of the encoder's hidden layers?


A: Let's base it on our decoder's current hidden state (our current representation of meaning) and all of the encoder's hidden layers!



seq2seq + Attention

Q: How do we determine how much to pay attention to each of the encoder's hidden layers?


A: Let's base it on our decoder's current hidden state (our current representation of meaning) and all of the encoder's hidden layers!



seq2seq + Attention

Q: How do we determine how much to pay attention to each of the encoder's hidden layers?


A: Let's base it on our decoder's current hidden state (our current representation of meaning) and all of the encoder's hidden layers!




seq2seq + Attention

Q: How do we determine how much to pay attention to each of the encoder's hidden layers?

A: Let's base it on our decoder's current hidden state (our current representation of meaning) and all of the encoder's hidden layers!



seq2seq + Attention



We multiply each encoder's hidden layer by its a_i^1 attention weights and sum all these vectors to create a context vector

$$c_1^D$$


$$h_1^D$$

Attention (softmax'd)

$$\begin{aligned}a_1^1 &= 0.51 \\a_2^1 &= 0.28 \\a_3^1 &= 0.14 \\a_4^1 &= 0.07\end{aligned}$$


seq2seq + Attention

REMEMBER: each attention weight a_i^j is based on the **decoder's** current hidden state, too.




seq2seq + Attention

REMEMBER: each attention weight a_i^j is based on the **decoder's** current hidden state, too.




seq2seq + Attention

REMEMBER: each attention weight a_i^j is based on the **decoder's** current hidden state, too.




seq2seq + Attention

REMEMBER: each attention weight a_i^j is based on the **decoder's** current hidden state, too.



seq2seq + Attention

REMEMBER: each attention weight a_i^j is based on the **decoder's** current hidden state, too.



For convenience, Attention calculation summarized on 1 slide




Photo credit: https://lena-voita.github.io/nlp_course/seq2seq_and_attention.html

Les principales fonctions ‘Score’

Name	Alignment score function	Citation
Content-base attention	$\text{score}(s_t, \mathbf{h}_i) = \text{cosine}[s_t, \mathbf{h}_i]$	Graves2014
Additive(*)	$\text{score}(s_t, \mathbf{h}_i) = \mathbf{v}_a^\top \tanh(\mathbf{W}_a[s_t; \mathbf{h}_i])$	Bahdanau2015
Location-Base	$\alpha_{t,i} = \text{softmax}(\mathbf{W}_a s_t)$ Note: This simplifies the softmax alignment to only depend on the target position.	Luong2015
General	$\text{score}(s_t, \mathbf{h}_i) = s_t^\top \mathbf{W}_a \mathbf{h}_i$ where \mathbf{W}_a is a trainable weight matrix in the attention layer.	Luong2015
Dot-Product	$\text{score}(s_t, \mathbf{h}_i) = s_t^\top \mathbf{h}_i$	Luong2015
Scaled Dot-Product(^)	$\text{score}(s_t, \mathbf{h}_i) = \frac{s_t^\top \mathbf{h}_i}{\sqrt{n}}$ Note: very similar to the dot-product attention except for a scaling factor; where n is the dimension of the source hidden state.	Vaswani2017

Il nous reste à construire le réseau

- ▶ Encodeur → inchangé
- ▶ Décodeur → il faut calculer l'attention

```
def build_decoder_attention():
    teacher_inputs = layers.Input(shape=(_OUTPUT_LENGTH,), name="decInput")
    ctx_input = [layers.Input(shape=(_RNN_SIZE,), name="ctxInputH"),
                layers.Input(shape=(_RNN_SIZE,), name="ctxInputC")]
    encoder_outputs = layers.Input(shape=(_INPUT_LENGTH, _RNN_SIZE))
```

```
h = layers.Embedding(input_dim=_VOCAB_SIZE, output_dim=_EMBEDDING_SIZE)(teacher_inputs)
decoder_outputs, memory_state, carry_state = layers.LSTM(_RNN_SIZE, return_sequences=True,
                                                       return_state=True,
                                                       dropout=_DROPOUT_RATE, recurrent_dropout=_DROPOUT_RATE)(h, initial_state=ctx_input)
decoder_context = [memory_state, carry_state]
```

```
attention_context, attention_scores = attention(decoder_outputs, encoder_outputs)
decoder_outputs = layers.concatenate([attention_context, decoder_outputs])
```



```
outputs = layers.Dense(_VOCAB_SIZE, activation='softmax')(decoder_outputs)
return models.Model([encoder_outputs, teacher_inputs, context_input],
                   [outputs, decoder_context, attention_scores], name="decoder")
```

- ▶ Seq2Seq → inchangé

Il nous reste à construire le réseau

- ▶ Encodeur → inchangé
- ▶ Decodeur
 - ▶ → il faut calculer l'attention
 - ▶ → il faut intégrer l'attention
- ▶ Solution I

Solution 2 : présentée dans les précédents transparents
On l'utilisera lors de TD (plus simple à impléter)



- ▶ Seq2Seq → inchangé

Calcul de l'attention

- ▶ Dépend évidemment de la fonction score

- ▶ def attention(query, value):

- ▶ Query → vient du decoder
 - ▶ Value → vient de l'encoder

- ▶ Par exemple pour 'dot-product'


```
attention = layers.dot([query, value], axes=[2,2], name="Att_score")
attention_scores = layers.Activation('softmax', name="Att_weight")(attention)
attention_context = layers.dot([attention_scores, value], axes=[2,1])
return attention_context, attention_scores
```

- ▶ Ou utilisation du layer Attention de Keras

```
attention_context, attention_scores = layers.Attention()(inputs=[query, value],
                                                       return_attention_scores=True)
```

Attention : ne fonctionne pas sur toutes les plate-formes, vous pouvez essayer un autre modèle d'attention disponible dans Keras

Afficher de la matrice d'attention



Possible uniquement
si on a les scores
d'attention

Lab of the day

- ▶ Tache : traduction
 - ▶ Des dates dans différents formats → date dans un seul format
 - ▶ On travaille au niveau des ‘caractères’
- ▶ Partie 1 : on construit le jeu de données
 - ▶ X, y et teacher
- ▶ Partie 2 : on construit, entraîne et prédit avec un modèle Seq2Seq
 - ▶ Encodeur
 - ▶ Decodeur
 - ▶ On assemble Encodeur et decodeur pour construire le modèle complet
 - ▶ On entraîne le réseau avec early stopping / On babysit
 - ▶ On prédit
- ▶ Partie 3 : Seq2Seq avec attention
 - ▶ On ajoute un calcul d’attention
 - ▶ On refait les mêmes étapes que précédemment
 - ▶ Enfin, on peut afficher la matrice d’attention pour un exemple donné
- ▶ **Les éléments pour construire le réseau de neurones, l’entraîner et prédire sont dans les transparents**

