

# Introduction to GPU Programming - 2

Fabrice Huet

# Take-away from Lab 1

- JIT Compiler
  - A kernel is compiled the first time it's run

```
fhuet@gpu ~/TP1> python3 TP1.py
0.33492976520210505  s
6.957538425922394e-05  s
4.2749568819999695e-05  s
3.902427852153778e-05  s
0.0002544168382883072  s
4.444550722837448e-05  s
```

- Either perform a dry run or exclude first measure

# Take-away from Lab 1

- JIT Compiler
    - Typing is inferred by the compiler
    - Can be checked with `kernel_name.inspect_llvm()`

```
{(int64, int64): ' ; ModuleID = "cuda.kernel.wrapper"\n\ttarget tri  
64:64:64-i1:8:8-i8:8:8-i16:16:16-i32:32:32-i64:64:64-i128:128:12  
8:128:128-n16:32:64"\n\ndeclare i32 @_ZN8__main__10computeIntB1  
EYakJSQB1PQBk0Bynm210iwU1a0UoLGHDPQE8oxrNQE_3dExx"(i8** %".1", i  
n__10computeIntB2v1B94cw51cXTLSUwv1sCUT9Ww0FEw09RRQPKzLTg4gaGKF  
x"(i64 %".1", i64 %".2")\n{\n.4:\n    %.5" = alloca i8*\n    store  
    %.8" = call i32 @_ZN8__main__10computeIntB2v1B94cw51cXTLSUwv1s  
0iwU1a0UoLGHDPQE8oxrNQE_3dExx"(i8** %".5", i64 %".1", i64 %".2")
```

- For benchmarking parameters
    - Write multiple kernels
    - Type parameters explicitly

```
@cuda.jit
def computeFloat(a,n) :
    r = 1.0
    for i in range(0, n):
        r*=a
```

```
def benchInt(runs):
    threadsPerBlock = 1024
    blocksPerGrid = 10
    result = []
    a = np.int32(2)
    n = 10000
    for i in range(runs):
        start = timer()
        computeInt[blocksPerGrid,blocksPerGrid](a,n)
        cuda.synchronize()
        dt = timer() - start
        print(" ", dt, " s")
        result.append(dt)
    print("Average Int 32 :", threadsPerBlock, np.average(result[1:]))
```

# Take-away from Lab 1

- Execution time is not the best metric
  - Need long running kernels to measure execution time
  - Increase threads => increase parallelism
    - So no increase in execution time...
- Prefer other metrics like ops/s or GB/s
  - Compute nb of op

```
@cuda.jit
def computeFloat(a,n) :
    r = 1.0
    for i in range(0, n):
        r*=a
        (n*threadsPerBlock*blocksPerGrid)
```

# Take-away from Lab 1

- Micro-benchmarks are hard!
  - Measurements vary widely from one run to another
  - Impact of other kernels running concurrently
  - Depends on the architecture (here RTX 3090)

```
Nb operations/s float64 : 2.06e+03 Gops/s  
Nb operations/s float32 : 2.72e+03 Gops/s  
Nb operations/s float16 : 2.14e+03 Gops/s  
Nb operations/s int32 : 6.96e+03 Gops/s
```

<b>Peak FP32 TFLOPS (non-Tensor)<sup>1</sup></b>	35.6
<b>Peak FP16 TFLOPS (non-Tensor)<sup>1</sup></b>	35.6
<b>Peak BF16 TFLOPS (non-Tensor)<sup>1</sup></b>	35.6
<b>Peak INT32 TOPS (non-Tensor)<sup>1,3</sup></b>	17.8

# Cuda & Numba

Device function

# Kernel and device functions

- A kernel is written in Python
  - Not all Python API is supported
- Add @cuda.jit in front of function

```
@cuda.jit  
def writeGlobalIDUnevenArray(array):
```

- A device function is similar
  - A device function can only be called **from** a kernel

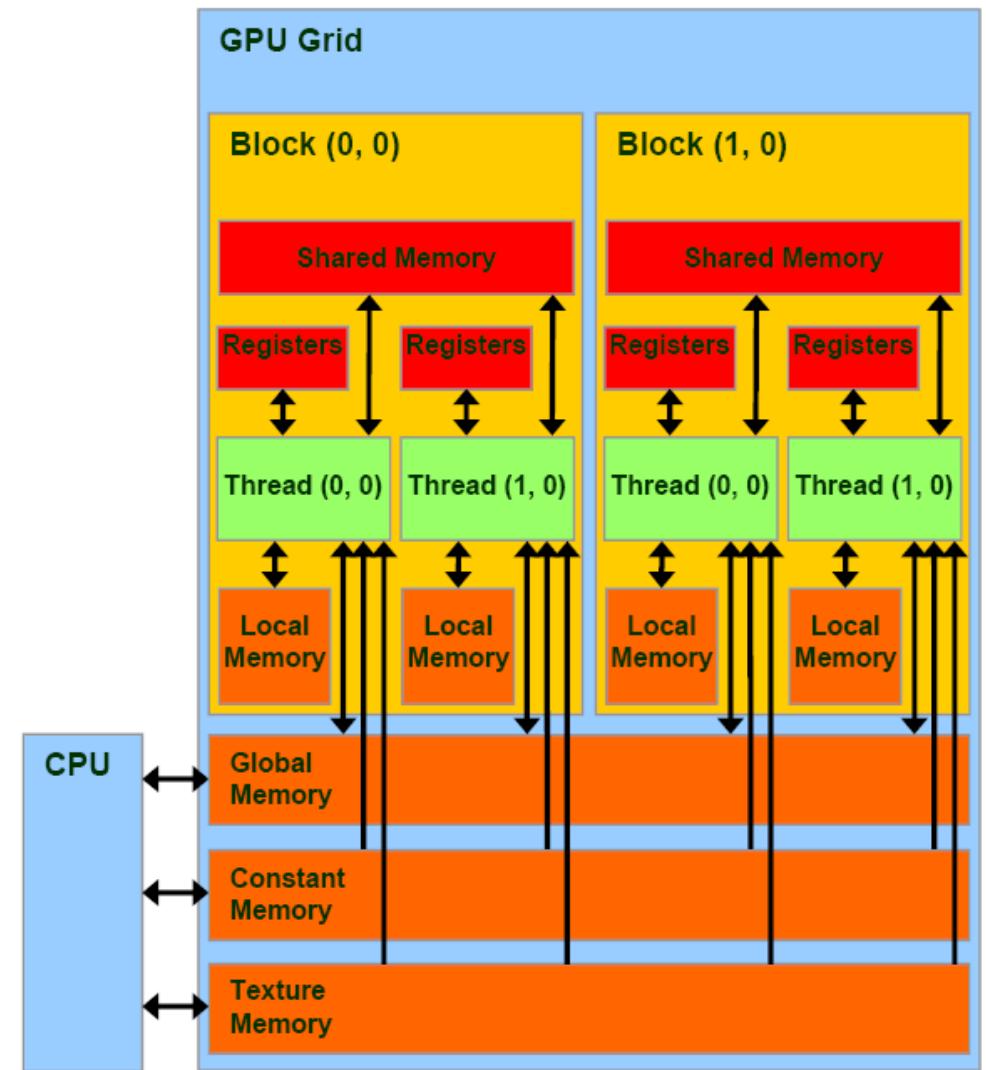
```
@cuda.jit(device=True)  
def deviceFunction(tab, index):
```

# Cuda & Numba

Memory organization

# NUMA

- Non Uniform Memory Access
- Host and device memory are separated
  - Explicit memory transferr
- Hierarchical memory
  - Different size, different access pattern, different performance



<https://www.3dgep.com/cuda-memory-model/>

# Memory

- Global Memory
  - Located on the device
  - Accessible from host or device
  - Used for transferring data between host and device
  - Limited by the amount of memory on the device
- Shared memory
  - Located on the device, roughly 100x faster than global
  - Only accessible from inside the device
  - Shared memory inside a SM
    - Sharing between threads of the same block
  - Very limited (64KB-100KB)
    - Shared with L1 cache

# Memory management

- Explicit memory management
  - Close to CUDA/C
  - Based on arrays
- Transferring an ndarray

```
numba.cuda.to_device(obj, stream=0, copy=True, to=None)
```

- Returns a DeviceNDArray
- DeviceNDArray
  - Reference to an ndarray on the device
  - Must be passed as argument to the kernel

```
#Generate array of size with zeros
A = np.zeros(size, dtype=np.uint16)
#Send array to device
d_A = cuda.to_device(A)
#Execute kernel
writeGlobalIDUnevenArray[ blocksPerGrid,threadsPerBlock](d_A)
```

# Memory management

- Creating an array directly on the device

```
numba.cuda.device_array(shape, dtype=np.float, strides=None, order='C', stream=0)
```

- By default use float
  - Type numpy
- Array\_like allocation

```
numba.cuda.device_array_like(ary, stream=0)
```

- Copy device to host (bring back the array)

```
copy_to_host(ary=None, stream=0)
```

# Memory management – Shared array

- Allocating shared array
  - Fast memory
  - Shared among all threads of the same block

```
numba.cuda.shared.array(shape, type)
```

- Must be inside a kernel or a device function
  - Shape must be known at compile time
    - Typically a global variable is used
  - Type is a numba type
    - Close to numpy types

# Memory management – Shared array

- Principle
  - Declare the array at beginning of kernel
  - Use threads to copy from global memory to shared array
  - Sync threads
  - Do some smart stuff ...
  - Copy back data from shared array to global memory
- Beware
  - Shared memory accessed using local ID
  - Global memory accessed using global ID

```
@cuda.jit
def kernel_with_shared(array_in, array_out):
    g_x = cuda.grid(1)
    s = cuda.shared.array(Thread_block, nb.int32)
    tx = cuda.threadIdx.x
    if g_x < array_in.shape[0]:
        s[tx] = array_in[g_x]
    cuda.syncthreads()
    #do some smart stuff here

    #smart stuff done, copy back the data
    if g_x < array_out.shape[0]:
        array_out[g_x] = s[tx]
```

# Cuda & Numba

Optimizing memory access

# Memory Coalescing

- CUDA tries to optimize access to global memory
  - Groups memory access from the same warp (32 threads)
  - Very efficient if consecutive memory addresses
- Ideally in a warp
  - Thread id 0 accesses memory N
  - Thread id 1 accesses memory N+1
  - ...
- Allows for DRAM burst (good)

# Example Matrix addition

- We want to have a kernel to perform
  - $D = A + B$
- Straightforward ....

```
@cuda.jit
def matrix_add_1(a, b, out):
    x, y = cuda.grid(2)
    out[x][y] = a[x][y] + b[x][y]

@cuda.jit
def matrix_add_2(a, b, out):
    x, y = cuda.grid(2)
    out[y][x] = a[y][x] + b[y][x]
```

- Is there one better ?
  - Depends on how matrix are organized in memory

## Row-major order

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$

## Column-major order

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$

# Example Matrix addition

- Numba relies on Numpy arrays
  - Numpy uses default C order
  - Row Major (<https://numpy.org/doc/1.14/glossary.html>)
- What about threads ?
  - Threads in a warp are linearized on Row-major also (x, y, z)
- Now let's go back to the kernels

```
@cuda.jit
def matrix_add_1(a, b, out):
    x, y = cuda.grid(2)
    out[x][y] = a[x][y] + b[x][y]

@cuda.jit
def matrix_add_2(a, b, out):
    x, y = cuda.grid(2)
    out[y][x] = a[y][x] + b[y][x]
```

# Example Matrix addition

- Assume square matrices of size > 32
  - 32 threads will run at once
  - At first all 32 threads will have  $x == 0$  and  $y$  variable
  - $a[x][y]$  accesses row  $x$  and column  $y$
- Matrix\_add\_1
  - Consecutive threads will access different row
  - No coalescing
- Matrix\_add\_2
  - Consecutive threads access different columns
    - Same row
  - Coalescing!
- In practice Matrix\_add\_2 is 2x faster

```
@cuda.jit
def matrix_add_1(a, b, out):
    x, y = cuda.grid(2)
    out[x][y] = a[x][y] + b[x][y]

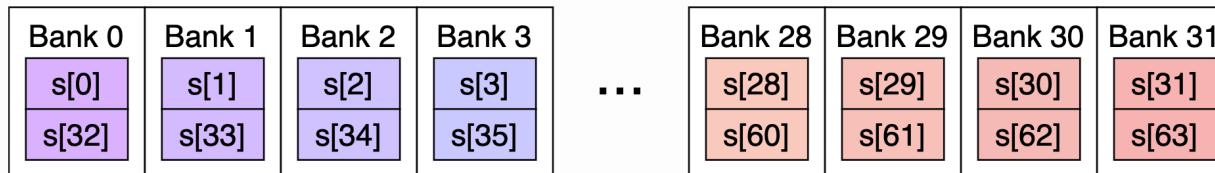
@cuda.jit
def matrix_add_2(a, b, out):
    x, y = cuda.grid(2)
    out[y][x] = a[y][x] + b[y][x]
```

# What about shared memory ?

- Shared memory is organized into 32 banks
  - 1 load/bank/clock cycle
  - Successive 32-bit words are stored in successive banks
- Banks can be accessed in parallel
  - So threads should avoid accessing the same shared memory address
- Otherwise bank conflict

```
_shared_ float s[64];
```

s[0]	s[1]	s[2]	s[3]	s[4]	s[5]	s[6]	s[7]	s[8]	s[9]	s[10]	s[11]	s[12]	s[13]	s[14]	s[15]	s[16]	s[17]	s[18]	s[19]	s[20]	s[21]	s[22]	s[23]	s[24]	s[25]	s[26]	s[27]	s[28]	s[29]	s[30]	s[31]
s[32]	s[33]	s[34]	s[35]	s[36]	s[37]	s[38]	s[39]	s[40]	s[41]	s[42]	s[43]	s[44]	s[45]	s[46]	s[47]	s[48]	s[49]	s[50]	s[51]	s[52]	s[53]	s[54]	s[55]	s[56]	s[57]	s[58]	s[59]	s[60]	s[61]	s[62]	s[63]



# Cheat sheets

Don't panic

# Before writing some code

- Identify which part will be executed on the GPU
  - This is your kernel
- Identify the data needed by the kernel
  - Size and shape
- Assume a “safe” thread block size to start
  - 32 usually good
- Deduce your grid size
  - Typically data size/thread block

# Writing some code

- Initialize/create your data on the host
- Transfer data to device
  - *cuda\_to\_device(...)*
- Call the kernel with the correct grid and thread block sizes
- Transfer data back to the host
  - *copy\_to\_host(...)*

# Writing a kernel

- You have to think like a thread
  - As a thread, what should you do ?
- Compute your global ID
  - Usually you need it
- Always check you don't read/write outside of your data
  - More on this during lab sessions
- Usually a for loop in an algorithm can be removed
  - Each index of the loop is handled by a thread
  - Data parallelism